



MACHINE LEARNING IN THE TIDYVERSE

Foundations of Tidy Machine Learning

Dmitriy (Dima) Gorenshteyn

Sr. Data Scientist,

Memorial Sloan Kettering Cancer Center



The Core of Tidy Machine Learning





The Core of Tidy Machine Learning





List Column Workflow

1

Make a
list column

```
nest()
```

2

Work with
list columns

```
map()
```

3

Simplify the
list columns

```
unnest()  
map_*()
```



The Gapminder Dataset

- **dslabs** package
- **Observations:** 77 countries for 52 years per country (1960-2011)
- **Features:**
 - year
 - infant_mortality
 - life_expectancy
 - fertility
 - population
 - gdpPercap



List Column Workflow

1

Make a
list column

nest()

2

Work with
list columns

`map()`

3

Simplify the
list columns

unnest()


`map_*()`

Step 1: Make a List Column - Nest Your Data

country	year	infant_mortality	life_expectancy	fertility	population	gdpPercap
Algeria	1960	148	47.5	7.65	11124892	1242
Algeria	1961	148	48	7.65	11404859	1047
Algeria	1962	148	48.6	7.65	11690152	820
Argentina	1960	59.9	65.4	3.11	20619075	5253
Argentina	1961	59.7	65.5	3.1	20953079	5450
Argentina	1962	59.6	65.6	3.09	21287682	5318
Australia	1960	20.3	70.9	3.45	10292328	9393
Australia	1961	20	71.1	3.55	10494911	9428
Australia	1962	19.5	70.9	3.43	10691220	9381
Austria	1960	37.3	68.8	2.7	7065525	7415
Austria	1961	35	69.7	2.79	7105654	7781
Austria	1962	32.9	69.5	2.8	7151077	7937

Step 1: Make a List Column - Nest Your Data

country	data	country	year	infant_mortality	life_expectancy	fertility	population	gdpPercap
Algeria	<tibble [52 x 66]>	Algeria	1960	148	47.5	7.65	11124892	1242
Argentina	<tibble [52 x 66]>	Algeria	1961	148	48	7.65	11404859	1047
Australia	<tibble [52 x 66]>	Algeria	1962	148	48.6	7.65	11690152	820
Austria	<tibble [52 x 66]>	Argentina	1960	59.9	65.4	3.11	20619075	5253
		Argentina	1961	59.7	65.5	3.1	20953079	5450
		Argentina	1962	59.6	65.6	3.09	21287682	5318
		Australia	1960	20.3	70.9	3.45	10292328	9393
		Australia	1961	20	71.1	3.55	10494911	9428
		Australia	1962	19.5	70.9	3.43	10691220	9381
		Austria	1960	37.3	68.8	2.7	7065525	7415
		Austria	1961	35	69.7	2.79	7105654	7781
		Austria	1962	32.9	69.5	2.8	7151077	7937



Nesting By Country

```
library(tidyverse)
nested <- gapminder %>%
  group_by(country) %>%
  nest()
```

country	data	country	year	infant_mortality	life_expectancy	fertility	population	gdpPercap
Algeria	<tibble [52 x 66]>	Algeria	1960	148	47.5	7.65	11124892	1242
Argentina	<tibble [52 x 66]>	Algeria	1961	148	48	7.65	11404859	1047
Australia	<tibble [52 x 66]>	Algeria	1962	148	48.6	7.65	11690152	820
Austria	<tibble [52 x 66]>	Argentina	1960	59.9	65.4	3.11	20619075	5253
		Argentina	1961	59.7	65.5	3.1	20953079	5450
		Argentina	1962	59.6	65.6	3.09	21287682	5318
		Australia	1960	20.3	70.9	3.45	10292328	9393
		Australia	1961	20	71.1	3.55	10494911	9428
		Australia	1962	19.5	70.9	3.43	10691220	9381
		Austria	1960	37.3	68.8	2.7	7065525	7415
		Austria	1961	35	69.7	2.79	7105654	7781
		Austria	1962	32.9	69.5	2.8	7151077	7937

Viewing a Nested Tibble

country	data		country	year	infant_mortality	life_expectancy	fertility	population	gdpPercap
Algeria	<tibble [52 x 66]>	nested\$data[[1]]	Algeria	1960	148	47.5	7.65	11124892	1242
Argentina	<tibble [52 x 66]>		Algeria	1961	148	48	7.65	11404859	1047
Australia	<tibble [52 x 66]>		Algeria	1962	148	48.6	7.65	11690152	820
Austria	<tibble [52 x 66]>	nested\$data[[4]]	Austria	1960	37.3	68.8	2.7	7065525	7415
			Austria	1961	35	69.7	2.79	7105654	7781
			Austria	1962	32.9	69.5	2.8	7151077	7937

Viewing a Nested Tibble

```
> nested$data[[4]]  
# A tibble: 52 x 6  
  year infant_mortality life_expectancy fertility population gdpPercap  
  <int>      <dbl>      <dbl>      <dbl>      <dbl>      <int>  
1  1960      37.3      68.8      2.70      7065525      7415  
2  1961      35.0      69.7      2.79      7105654      7781  
3  1962      32.9      69.5      2.80      7151077      7937  
4  1963      31.2      69.6      2.82      7199962      8209  
5  1964      29.7      70.1      2.80      7249855      8652  
6  1965      28.3      69.9      2.70      7298794      8893
```

Step 3: Simplify List Columns - unnest()

country	data	country	year	infant_mortality	life_expectancy	fertility	population	gdpPercap
Algeria	<tibble [52 x 66]>	Algeria	1960	148	47.5	7.65	11124892	1242
Argentina	<tibble [52 x 66]>	Algeria	1961	148	48	7.65	11404859	1047
Australia	<tibble [52 x 66]>	Algeria	1962	148	48.6	7.65	11690152	820
Austria	<tibble [52 x 66]>	Argentina	1960	59.9	65.4	3.11	20619075	5253
		Argentina	1961	59.7	65.5	3.1	20953079	5450
		Argentina	1962	59.6	65.6	3.09	21287682	5318
		Australia	1960	20.3	70.9	3.45	10292328	9393
		Australia	1961	20	71.1	3.55	10494911	9428
		Australia	1962	19.5	70.9	3.43	10691220	9381
		Austria	1960	37.3	68.8	2.7	7065525	7415
		Austria	1961	35	69.7	2.79	7105654	7781
		Austria	1962	32.9	69.5	2.8	7151077	7937



unnest()

Step 3: Simplify List Columns - unnest()

```
nested %>%  
  unnest(data)  
  
# A tibble: 4,004 x 7  
  country    year infant_mortality life_expectancy fertility population ...  
  <fct>    <int>         <dbl>         <dbl>         <dbl>         <dbl>    ...  
1 Algeria  1960           148           47.5           7.65      11124892    ...  
2 Algeria  1961           148           48.0           7.65      11404859    ...  
3 Algeria  1962           148           48.6           7.65      11690152    ...  
4 Algeria  1963           148           49.1           7.65      11985130    ...  
5 Algeria  1964           149           49.6           7.65      12295973    ...  
6 Algeria  1965           149           50.1           7.66      12626953    ...
```



MACHINE LEARNING IN THE TIDYVERSE

Let's Get Started!



MACHINE LEARNING IN THE TIDYVERSE

The map family of functions

Dmitriy (Dima) Gorenshteyn

Sr. Data Scientist,

Memorial Sloan Kettering Cancer Center



List Column Workflow

1

Make a
list column

```
nest()
```

2

Work with
list columns

```
map()
```

3

Simplify the
list columns

```
unnest()  
map_*()
```




List Column Workflow

1

Make a
list column

`nest()`

2

Work with
list columns

`map()`

3

Simplify the
list columns

`unnest()``map_*()`




The map Function

`map(.x = , .f =)`

The map Function

map(.x = , .f =)




`.x = [vector]`
or
`.x = [[list]]`

`.f = function()`
or
`.f = ~formula`

The map Function

map(.x = , .f =)



`.x = [vector]`
or
`.x = [[list]]`

`.f = mean`
or
`.f = ~mean(.x)`

Population Mean by Country

country	data
Algeria	<tibble [52 x 66]>
Argentina	<tibble [52 x 66]>
Australia	<tibble [52 x 66]>
Austria	<tibble [52 x 66]>

nested\$data[[1]]

year	infant_mortality	life_expectancy	fertility	population	gdpPercap
1960	148	47.5	7.65	11124892	1242
1961	148	48	7.65	11404859	1047
1962	148	48.6	7.65	11690152	820

```
mean(nested$data[[1]]$population)
[1] 23129438
```



Population Mean by Country

```
map(.x = nested$data, .f = ~mean(.x$population))
```

```
[[1]]  
[1] 23129438
```

```
[[2]]  
[1] 30783053
```

```
[[3]]  
[1] 16074837
```

```
[[4]]  
[1] 7746272
```

2: Work with List Columns - map() and mutate()

```
pop_df <- nested %>%  
  mutate(pop_mean = map(data, ~mean(.x$population)))
```

```
pop_df  
  
# A tibble: 77 x 3  
  country      data      pop_mean  
  <fct>      <list>      <list>  
1 Algeria  <tibble [52 x 6]> <dbl [1]>  
2 Argentina <tibble [52 x 6]> <dbl [1]>  
3 Australia <tibble [52 x 6]> <dbl [1]>  
4 Austria   <tibble [52 x 6]> <dbl [1]>  
5 Bangladesh <tibble [52 x 6]> <dbl [1]>  
6 Belgium   <tibble [52 x 6]> <dbl [1]>
```

3: Simplify List Columns - unnest()

```
pop_df %>%  
  unnest(pop_mean)  
  
# A tibble: 77 x 3  
  country      data      pop_mean  
  <fct>      <list>      <dbl>  
1 Algeria    <tibble [52 x 6]> 23129438  
2 Argentina  <tibble [52 x 6]> 30783053  
3 Australia  <tibble [52 x 6]> 16074837  
4 Austria    <tibble [52 x 6]>  7746272  
5 Bangladesh <tibble [52 x 6]> 97649407  
6 Belgium    <tibble [52 x 6]>  9983596
```


List Column Workflow

1

Make a
list column

```
group_by(gapminder, country) %>%  
  nest() %>%
```

2

Work with
list columns

```
  mutate(pop_mean =  
    map(data, ~mean(.x$population)) %>%
```

3

Simplify the
list columns

```
    unnest(pop_mean)
```



Work With + Simplify List Columns With `map_*()`

function	returns
<code>map()</code>	list
<code>map_dbl()</code>	double
<code>map_lgl()</code>	logical
<code>map_chr()</code>	character
<code>map_int()</code>	integer

Work With + Simplify List Columns With map_dbl()

```
nested %>%  
  mutate(pop_mean = map_dbl(data, ~mean(.x$population)))  
  
# A tibble: 77 x 3  
  country      data      pop_mean  
  <fct>      <list>      <dbl>  
1 Algeria    <tibble [52 x 6]> 23129438  
2 Argentina  <tibble [52 x 6]> 30783053  
3 Australia  <tibble [52 x 6]> 16074837  
4 Austria    <tibble [52 x 6]>  7746272  
5 Bangladesh <tibble [52 x 6]> 97649407  
6 Belgium    <tibble [52 x 6]>  9983596
```

Build Models with map()

```
nested %>%
  mutate(model = map(data, ~lm(formula = population~fertility, data = .x)))

# A tibble: 77 x 3
  country      data      model
  <fct>      <list>    <list>
1 Algeria <tibble [52 x 6]> <S3: lm>
2 Argentina <tibble [52 x 6]> <S3: lm>
3 Australia <tibble [52 x 6]> <S3: lm>
4 Austria <tibble [52 x 6]> <S3: lm>
5 Bangladesh <tibble [52 x 6]> <S3: lm>
6 Belgium <tibble [52 x 6]> <S3: lm>
```



MACHINE LEARNING IN THE TIDYVERSE

Let's map something!



MACHINE LEARNING IN THE TIDYVERSE

Tidy your models with broom

Dmitriy (Dima) Gorenshteyn

Sr. Data Scientist,

Memorial Sloan Kettering Cancer Center

List Column Workflow

1

Make a
`list column`

2

Work with
`list columns`

3

Simplify the
`list columns`

```
library(broom)
library(yardstick)
library(rsample)
...
```

List Column Workflow

1

Make a
`list column`

2

Work with
`list columns`

3

Simplify the
`list columns`

```
library(broom)  
library(yardstick)  
library(rsample)  
...
```




Broom Toolkit

- **tidy()**: returns the statistical findings of the model (such as coefficients)
- **glance()**: returns a concise one-row summary of the model
- **augment()**: adds prediction columns to the data being modeled

Summary of algeria_model

```
> summary(algeria_model)
```

Call:

```
lm(formula = life_expectancy ~ year, data = .x)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.044	-1.577	-0.543	1.700	3.843

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.197e+03	3.994e+01	-29.96	<2e-16 ***
year	6.349e-01	2.011e-02	31.56	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.177 on 50 degrees of freedom

Multiple R-squared: 0.9522, Adjusted R-squared: 0.9513

F-statistic: 996.2 on 1 and 50 DF, p-value: < 2.2e-16



tidy()

```
> summary(algeria_model)
```

Call:

```
lm(formula = life_expectancy ~ year, data = .x)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.044	-1.577	-0.543	1.700	3.843

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.197e+03	3.994e+01	-29.96	<2e-16 ***
year	6.349e-01	2.011e-02	31.56	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.177 on 50 degrees of freedom

Multiple R-squared: 0.9522, Adjusted R-squared: 0.9513

F-statistic: 996.2 on 1 and 50 DF, p-value: < 2.2e-16



tidy()

```
library(broom)
```

```
tidy(algeria_model)
```

	term	estimate	std.error	statistic	p.value
1	(Intercept)	-1196.5647772	39.93891866	-29.95987	1.319126e-33
2	year	0.6348625	0.02011472	31.56209	1.108517e-34

glance()

```
> summary(algeria_model)
```

Call:

```
lm(formula = life_expectancy ~ year, data = .x)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.044	-1.577	-0.543	1.700	3.843

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.197e+03	3.994e+01	-29.96	<2e-16 ***
year	6.349e-01	2.011e-02	31.56	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.177 on 50 degrees of freedom

Multiple R-squared: 0.9522, Adjusted R-squared: 0.9513

F-statistic: 996.2 on 1 and 50 DF, p-value: < 2.2e-16



glance()

```
glance(algeria_model)
```

r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik
0.9522064	0.9512505	2.176948	996.1653	1.108517e-34	2	-113.2171
AIC	BIC	deviance	df.residual			
232.4342	238.288	236.9552	50			



augment()

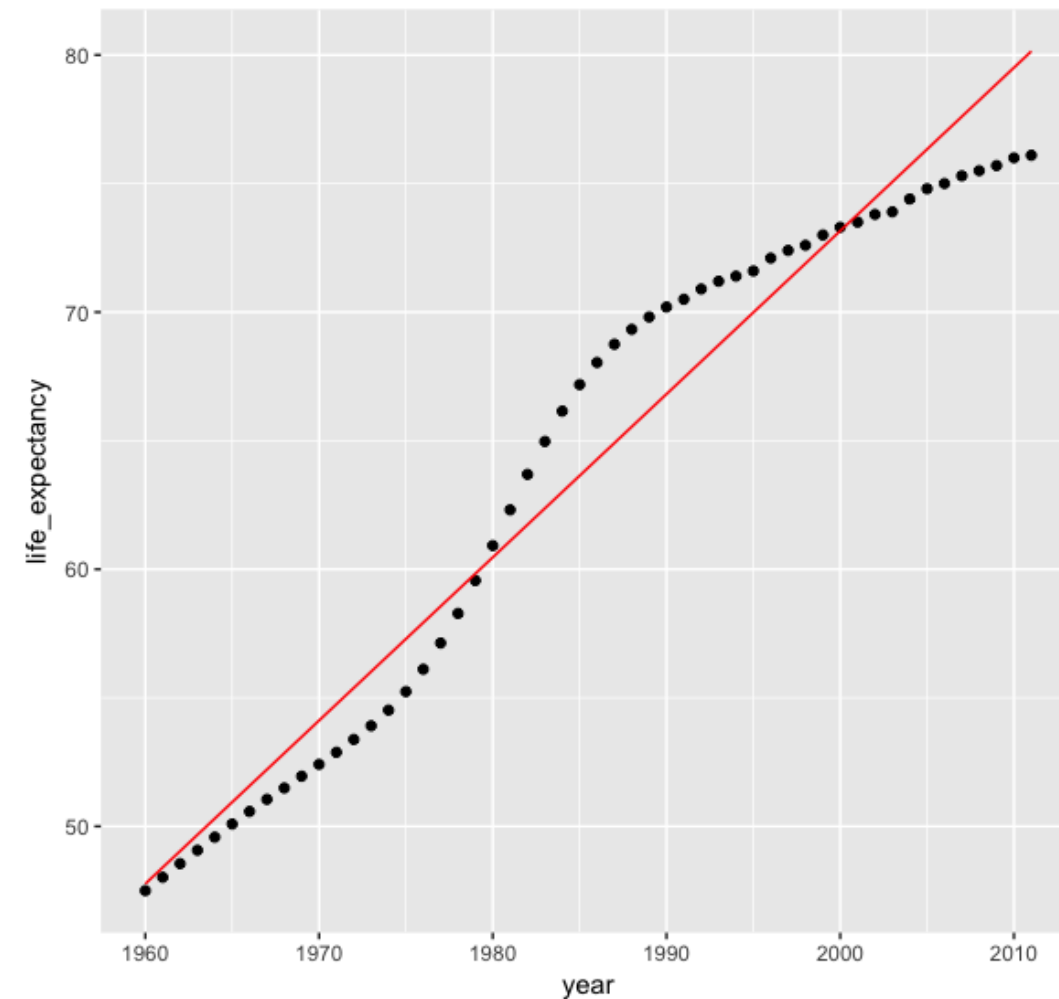
```
augment(algeria_model)
```

	life_expectancy	year	.fitted	.se.fit	.resid	.hat	.sigma
1	47.50	1960	47.76581	0.5951714	-0.2658128	0.07474601	2.198695
2	48.02	1961	48.40068	0.5779264	-0.3806753	0.07047725	2.198326
3	48.55	1962	49.03554	0.5608726	-0.4855379	0.06637924	2.197878
4	49.07	1963	49.67040	0.5440279	-0.6004004	0.06245198	2.197265
5	49.58	1964	50.30526	0.5274124	-0.7252630	0.05869547	2.196455
6	50.09	1965	50.94013	0.5110485	-0.8501255	0.05510971	2.195498



Plotting Augmented Data

```
augment(algeria_model) %>%  
  ggplot(mapping = aes(x = year)) +  
  geom_point(mapping = aes(y = life_expectancy)) +  
  geom_line(mapping = aes(y = .fitted), color = "red")
```





MACHINE LEARNING IN THE TIDYVERSE

Let's use broom!