





SUPERVISED LEARNING IN R: CLASSIFICATION

Making binary predictions with regression

Brett Lantz
Instructor




Introducing linear regression






Regression for binary classification




Introducing logistic regression






Making predictions with logistic regression



```
m <- glm(y ~ x1 + x2 + x3,  
         data = my_dataset,  
         family = "binomial")
```



```
prob <- predict(m, test_dataset,  
               type = "response")
```

```
pred <- ifelse(prob > 0.50, 1, 0)
```



SUPERVISED LEARNING IN R: CLASSIFICATION

Let's practice!




SUPERVISED LEARNING IN R: CLASSIFICATION

Model performance tradeoffs




Brett Lantz
Instructor

Understanding ROC curves






Area under the ROC curve





Using AUC and ROC appropriately





SUPERVISED LEARNING IN R: CLASSIFICATION

Let's practice!



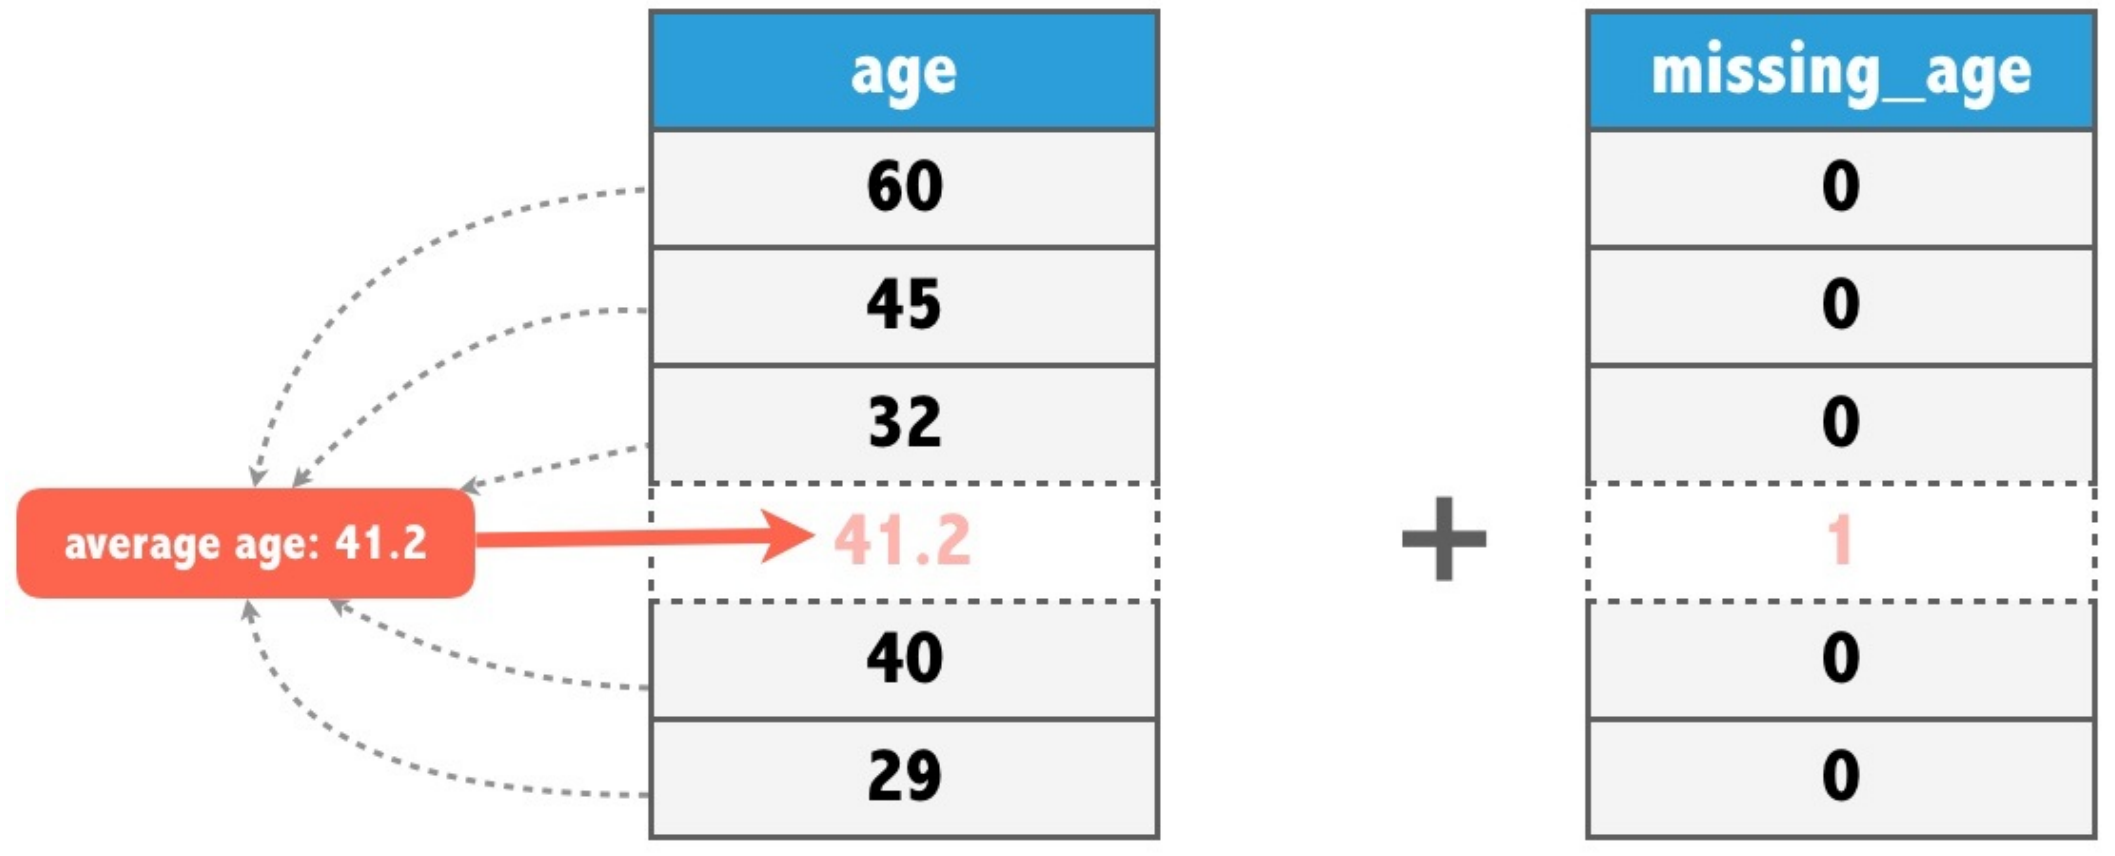
SUPERVISED LEARNING IN R: CLASSIFICATION

Dummy variables, missing data, and interactions

Brett Lantz
Instructor




Imputing missing data





Interaction effects



```
# interaction of obesity and smoking  
glm(disease ~ obesity * smoking,  
    data = health,  
    family = "binomial")
```



SUPERVISED LEARNING IN R: CLASSIFICATION

Let's practice!



SUPERVISED LEARNING IN R: CLASSIFICATION




Automatic feature selection

Brett Lantz
Instructor





Stepwise regression


**backward
stepwise**



**forward
stepwise**



Stepwise regression caveats





SUPERVISED LEARNING IN R: CLASSIFICATION

Let's practice!